
DreamDiffusion

A Preprint

1 Main

3D gen models can be trained on voxels or point clouds, but data is relatively scarce compared to plentiful 2D images.

2 Introduction

Text-conditioned diffusion models achieve 2D synthesis by learning to reverse a noising process, but they do not directly produce 3D geometry. DreamFusion overcomes this by using a frozen 2D diffusion network as a loss to sculpt a Neural Radiance Field (NeRF) to match a text prompt, without any paired 3D data.

3 Diffusion Models and Score Distillation Sampling

3.1 2D Text-to-Image Diffusion

A diffusion model defines a forward process

$$q(z_t | x) = \mathcal{N}(\alpha_t x, \sigma_t^2 I),$$

and learns a denoiser ϵ_ϕ via

$$\mathcal{L}_{\text{Diff}}(\phi, x) = \mathbb{E}_{t, \epsilon} [w(t) \|\epsilon_\phi(\alpha_t x + \sigma_t \epsilon; t, y) - \epsilon\|^2]. \quad (1)$$

3.2 Score Distillation Sampling

We wish to optimize generator parameters θ so that $x = g(\theta)$ lies in a high-likelihood region of the diffusion model. Define the SDS loss gradient as

$$\nabla_\theta \mathcal{L}_{\text{SDS}} = \mathbb{E}_{t, \epsilon} \left[w(t) (\hat{\epsilon}_\phi(z_t; y, t) - \epsilon) \frac{\partial g(\theta)}{\partial \theta} \right], \quad z_t = \alpha_t g(\theta) + \sigma_t \epsilon, \quad (2)$$

omitting the expensive U-Net Jacobian and absorbing constant factors into $w(t)$.

3.3 Sampling in Parameter Space

Existing diffusion sampling operates in pixel space, matching the data the model was trained on. Instead, we parameterize images via a differentiable generator $x = g(\theta)$ (e.g. a NeRF) and seek

$$\theta^* = \arg \min_{\theta} \mathcal{L}_{\text{Diff}}(\phi, g(\theta)),$$

but found this unstable in practice. By dropping the U-Net Jacobian term, we obtain the practical update of Eq. (2), which follows the learned score function to guide $g(\theta)$ toward regions the diffusion model deems plausible.

4 Neural Rendering of a 3D Model

A NeRF MLP maps each point $\mu \in \mathbb{R}^3$ to density τ and color ρ :

$$(\tau, \rho) = \text{MLP}(\mu; \theta).$$

Colors accumulate along rays via

$$C = \sum_i w_i c_i, \quad w_i = \alpha_i \prod_{j < i} (1 - \alpha_j), \quad \alpha_i = 1 - \exp(-\tau_i \delta_i),$$

and each sample’s shaded color is

$$c_i = \rho_i \circ (\ell_\rho \max(0, n_i \cdot \frac{\ell - \mu_i}{\|\ell - \mu_i\|}) + \ell_a),$$

with normals $n_i = -\nabla_\mu \tau_i / \|\nabla_\mu \tau_i\|$.

5 Text-to-3D Optimization

1. Initialize $\theta \sim \mathcal{N}(0, I)$.
2. Repeat for $k = 1, \dots, N$:
 - (a) Sample a random camera and lighting.
 - (b) Render $x = g(\theta)$ at low resolution.
 - (c) Draw $t \sim U[0, 1]$, $\epsilon \sim \mathcal{N}(0, I)$.
 - (d) Compute $z_t = \alpha_t x + \sigma_t \epsilon$.
 - (e) Query diffusion model: $\hat{\epsilon} = \epsilon_\phi(z_t; y, t)$.
 - (f) Compute gradient $\nabla_\theta \mathcal{L}_{\text{SDS}}$ via Eq. (2).
 - (g) Update $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}_{\text{SDS}}$.

6 Algorithm (Pseudocode)

Algorithm 1 DreamFusion Optimization Loop

Input: text prompt y , frozen diffusion ϵ_ϕ , renderer $g(\theta)$ Initialize $\theta \sim \mathcal{N}(0, I)$ $k = 1$ N Sample camera pose, light source $x \leftarrow g(\theta)$ $t \sim U(0, 1)$, $\epsilon \sim \mathcal{N}(0, I)$ $z_t \leftarrow \alpha_t x + \sigma_t \epsilon$ $\hat{\epsilon} \leftarrow \epsilon_\phi(z_t; y, t)$ $g_\theta \leftarrow w(t) (\hat{\epsilon} - \epsilon) \nabla_\theta x$ $\theta \leftarrow \theta - \eta g_\theta$ return θ

A Appendix A.4: Equivalence to Density Distillation

Here we show that the SDS gradient arises from a weighted KL divergence between the true noising distribution and the diffusion model’s predicted distribution. Define

$$\mathcal{L}_{\text{SDS}}(\phi, \theta) \triangleq \mathbb{E}_t \left[\frac{\sigma_t}{\alpha_t} w(t) \text{KL}(q(z_t | g(\theta); y, t) \| p_\phi(z_t; y, t)) \right].$$

Then one can verify

$$\nabla_\theta \mathcal{L}_{\text{SDS}} = \nabla_\theta \mathbb{E}_{t, \epsilon} \left[\frac{\sigma_t}{\alpha_t} w(t) \left\langle \nabla_{z_t} \log q(z_t | x), \frac{\partial g(\theta)}{\partial \theta} \right\rangle \right],$$

and using the score-matching identity $\nabla_{z_t} \log q(z_t | x) = -(\hat{\epsilon}_\phi(z_t; y, t) - \epsilon)/\sigma_t$, this yields exactly the practical gradient of Eq. (2) up to constant factors.