

Flash Attention 2

A Preprint

In this paper authors work on parallelism and work partitioning to address the problem or reaching only 30-40% of theoretical maximum

$Q, K, V \in \mathbb{R}^{N \times d}$ where N is length of input sequence and d is head dimension; $O \in \mathbb{R}^{N \times d}$ output. $S = QK^T$ $P = \text{softmax}(S)$ $O = PV$

Then backprop:

$$dV = P^T dO \quad (1)$$

$$dP = dOV^T \quad (2)$$

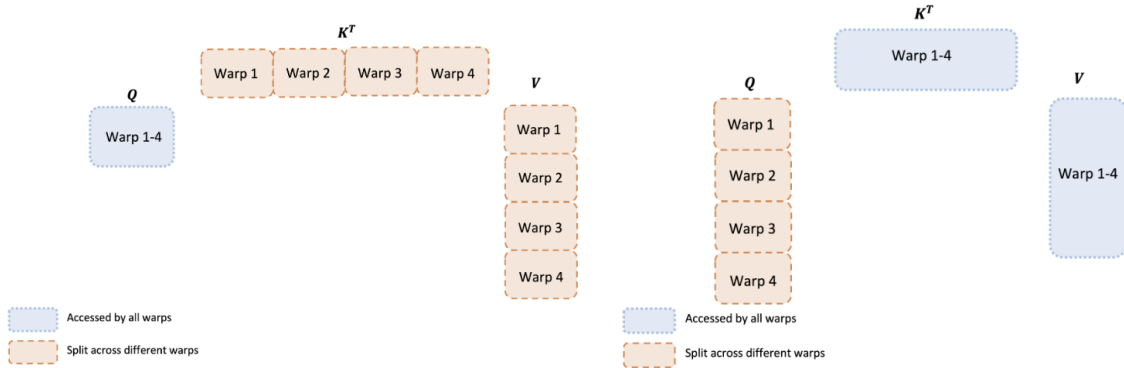
$$dS = d\text{softmax}(dP) \quad (3)$$

$$dQ = dSK \quad (4)$$

$$dK = QdS^T \quad (5)$$

Flash Attention 2 presents new ideas above first paper. More specifically authors elaborate on removing scaling term and saving statistics. As for removing scaling term $\text{diag}(l^{(2)})^{-1}$ is mentioned. In every pass while calculating output it was done next way: $O^{(2)} = \text{diag}(l^{(1)}/l^{(2)})^{-1}O^{(1)} + \text{diag}(l^{(2)})^{-1}e^{S^{(2)}-m^{(2)}}V^{(2)}$. The new way is presented: maintaining unscaled version of calculations and scaling only in the end of the loop: $O^{(2)} = \text{diag}(l^{(1)})^{-1}O^{(1)} + e^{S^{(2)}-m^{(2)}}V^{(2)}$. Also authors mention that it is not necessary to store max m and sum of exponentials l as logsumexp can be stored: $L^{(j)} = m^{(j)} + \log(l^{(j)})$

Authors also work on parallelization suggesting splitting rows into blocks which results in GPU load. Also instead of splitting K and V authors split Q across 4 warps while keeping K and V accessible by all warps. After each warp performs matrix multiply to get a slice of QK^T , they just need to multiply with their shared slice of V to get their corresponding slice of the output.



(a) FLASHATTENTION

(b) FLASHATTENTION-2

Algorithm 1 FlashAttention-2 forward pass

Require: Matrices $Q, K, V \in \mathbb{R}^{N \times d}$ in HBM, block sizes B_c, B_r .

- 1: Divide Q into $T_r = \lceil N/B_r \rceil$ blocks Q_1, \dots, Q_{T_r} of size $B_r \times d$, and divide K, V into $T_c = \lceil N/B_c \rceil$ blocks $K_1, \dots, K_{T_c}, V_1, \dots, V_{T_c}$ of size $B_c \times d$.
 - 2: Divide output $O \in \mathbb{R}^{N \times d}$ into T_r blocks O_1, \dots, O_{T_r} of size $B_r \times d$, and logsumexp L into T_r blocks L_1, \dots, L_{T_r} of size B_r .
 - 3: for $i = 1, \dots, T_r$ do
 - 4: Load Q_i from HBM to on-chip SRAM.
 - 5: Initialize

$$O_i^{(0)} = 0^{B_r \times d}, \quad \ell_i^{(0)} = 0^{B_r}, \quad m_i^{(0)} = (-\infty)^{B_r}.$$
 - 6: for $j = 1, \dots, T_c$ do
 - 7: Load K_j, V_j from HBM to on-chip SRAM.
 - 8: $S_i^{(j)} = Q_i K_j^\top \in \mathbb{R}^{B_r \times B_c}$.
 - 9: $m_i^{(j)} = \max(m_i^{(j-1)}, \text{rowmax}(S_i^{(j)}))$.
 - 10: $\tilde{P}_i^{(j)} = \exp(S_i^{(j)} - m_i^{(j)}) \in \mathbb{R}^{B_r \times B_c}$.
 - 11: $\ell_i^{(j)} = e^{m_i^{(j-1)} - m_i^{(j)}} \ell_i^{(j-1)} + \text{rowsum}(\tilde{P}_i^{(j)})$.
 - 12: $O_i^{(j)} = e^{m_i^{(j-1)} - m_i^{(j)}} O_i^{(j-1)} + \tilde{P}_i^{(j)} V_j$.
 - 13: end for
 - 14: $O_i = \text{diag}(\ell_i^{(T_c)})^{-1} O_i^{(T_c)}$.
 - 15: $L_i = m_i^{(T_c)} + \log(\ell_i^{(T_c)})$.
 - 16: Write O_i and L_i back to HBM.
 - 17: end for
 - 18: return O, L
-